

# Applications of Logistic Regression in Protein Contact Map Prediction

**Sritika Chakladar<sup>1, \*, +</sup>, Rahul Alapati<sup>1, +</sup>, Vineet Nayak<sup>1, +</sup>, Anuj Gupta<sup>1, +</sup> and Austin Ream<sup>2, +</sup>**

<sup>1</sup> Graduate Student, Computer Science and Software Engineering, Auburn University, 36830, USA <sup>2</sup> Undergraduate Student, Computer Science and Software Engineering, Auburn University, 36830, USA

\* szc0098@auburn.edu

+ these authors contributed equally to this work

## ABSTRACT

Prediction of protein contact map is of greater importance since it can improve the prediction of protein 3D structure. Protein contacts contain key information for the understanding of protein structure and function. This report lays out the implementation of Logistic Regression a non-linear regression algorithm, that has been used to train a binary logistic model to estimate the probability of contacts based on the position specific scoring matrices generated by PSI-BLAST. Two variants of gradient ascent namely: Batch & Stochastic and MCLE estimation have been used to learn the weight vectors during model training. Accuracies of top L/10, L/5 and L/2 predicted contacts are calculated to evaluate the classification performance of the model.

## 1. Introduction

To understand the mechanism of protein functions at the molecular level, it is usually necessary to determine their structures. However, determining 3D protein structure by experimental methods are high-cost, time-consuming and challenging. Fortunately, protein contact map or residue-residue contact map prediction simplifies the problem of protein structure prediction. The contacts in a protein sequence are predicted in the form a contact map, and then used in 3D protein structure prediction.

Contact Map of a protein, is a binary symmetric matrix and is a useful and simplified representation of the 3D protein structure. Contact Map contains important structural information of a protein, especially the backbone information and hence is crucial in predicting the 3D protein structure. The protein contact map prediction refers to the problem of predicting the spatial closeness of residue pairs in a folded protein structure. Contacts occurring between

sequentially distant residues, i.e. long-range contacts, impose strong constraints on the 3D structure of a protein and are particularly important for structural analysis, understanding the folding process and predicting the 3D structure. Even a small set of correctly predicted long-range contacts can be useful for improving ab initio structure prediction for proteins without known templates.

Logistic regression is a classification algorithm which is used to assign observations to a discrete set of classes. It transforms the input training dataset using a logistic function to return a probability which is mapped to discrete response classes based upon the prediction threshold. The learning algorithm uses regression coefficient to reduce the prediction error. It discovers the best value of the coefficient based on the prediction error and iteratively updates them until the stopping criteria is met.

Gradient ascent optimization is an iterative algorithm for finding the maximum of the cost function by taking steps proportional to the positive of the gradient at the current point. It has two variants: Batch and Stochastic Gradient Ascent. Batch gradient ascent calculates the gradient of the cost function by using the sum of the cost of each sample in every iteration. On the other hand, stochastic gradient ascent is a stochastic approximation of the gradient ascent optimization algorithm and it's an iterative method for maximizing the cost function over a single instance in the dataset. We use the gradient ascent optimization algorithm to update the regression coefficient vector to reach convergence quickly.

In this project, we train a binary logistic model to estimate the probability of contacts based on the position specific scoring matrices generated by PSI-BLAST. We implement both the variants of gradient ascent namely: Batch & Stochastic and MCLE estimation to learn the weight vectors during model training. We use these models to predict the residue-residue contacts in a protein sequence.

## **2. Methods**

The protein sequences in Multi-FASTA format, the actual intra residue distances (d) from the rr file and the position specific scoring matrices have been used to train and test the binary logistic model. Gradient ascent optimization algorithms and MCLE Estimation have been used to learn the weight vector. The learned weight vector is used in residue-residue contact prediction in a protein sequence.

The input protein sequences and their respective actual intra residue distances (d) classes have been curated into non-overlapping sets of Training (75%) and

Test (25%) datasets using the simple random sampling without replacement.

Now, PSIBLAST and the nr database have been used to generate train and test PSSMs for protein sequences in train and test datasets, respectively.

BLAST (Basic Local Alignment Search Tool) is a sequence similarity search method, in which a query protein is compared to protein sequences in a target database to identify regions of local alignment and report those alignments that score above a given score threshold. Position-Specific Iterative (PSI)-BLAST is a protein sequence profile search method that builds off the alignments generated by a run of the BLASTp program.

A sample PSSM for a sequence is shown below in figure 1:

>sequence

TIKVLFDVDDHEMVRIGISSYLSTQSDIEVVGEGASGKEAIAKAHELKPDILMDLLMEDMDGVEATT  
QIKKDLPQIKVLMSTFIEDKEVYRALDAGVDSYILKTTSAKDIADAVRKTSRGESVFEPEVLVKMR  
NRMKKRAELYEMLTEREMEILLIAKGYSNQEIASASHITIKTVKTHVSNILSKLEVQDRTQAVIYAF  
QHNLIQ

	A	R	N	D	C	Q	E	G	H	I	L	K	M	F	P	S	T	W	Y	V	
1 T	-1	-3	-2	-3	-3	-2	-2	-3	-3	-3	-3	-3	-2	-3	-4	-3	3	6	-4	-3	-2
2 I	-3	-5	-6	-6	-3	-5	-5	-6	-6	7	0	-5	2	-2	-5	-5	-2	-5	-4	2	
3 K	-3	7	-2	-3	-5	0	-2	-4	-3	-5	-5	4	-4	-5	-4	-1	1	-5	-4	-4	
4 V	-2	-5	-5	-6	-3	-5	-5	-6	-6	3	1	-5	-1	-3	-5	-4	-2	-5	-4	6	
5 L	-2	-5	-6	-6	-3	-4	-5	-6	-5	1	5	-5	3	-1	-5	-4	-3	-4	-3	2	
6 F	-3	-5	-6	-6	-3	-5	-5	-6	-5	5	4	-5	-1	1	-5	-5	-3	-4	-3	3	
7 V	4	-4	-5	-5	1	-4	-4	-4	-5	0	-2	-4	-2	-4	-4	-3	-2	-5	-4	6	
8 D	-4	-4	-1	8	-6	-2	0	-4	-3	-6	-6	-3	-6	-6	-4	-2	-3	-7	-6	-6	
9 D	-4	-4	-1	8	-6	-3	-1	-4	-3	-6	-6	-3	-6	-6	-4	-2	-3	-7	-6	-6	
10 H	-3	-2	-2	-3	-5	6	-1	-4	9	-5	-5	-2	-3	-4	-4	-3	-3	-5	-1	-5	
11 E	3	-1	-2	0	-4	1	4	-2	3	-4	-3	-2	0	-4	2	0	-1	-5	-4	-2	
12 M	-3	-4	-5	-6	-4	-4	-5	-6	-5	2	4	-4	7	-2	-5	-4	-3	-4	-3	3	
13 V	-2	-5	-5	-6	-3	-5	-5	-6	-5	2	1	-5	1	0	-5	-4	-3	-5	-3	6	
14 R	-4	8	-3	-4	-6	-1	-2	-5	-3	-5	-3	0	-4	-5	-5	-3	-3	-5	-4	-5	
15 I	1	2	-3	-1	-4	3	0	-2	-3	1	0	1	3	-4	-4	-1	1	-5	-4	-2	
16 G	-1	-5	-3	-4	-5	-4	-4	7	-4	-6	-6	-4	-5	-6	-4	-2	-4	-5	-5	-6	
17 I	-4	-5	-6	-6	-4	-5	-5	-6	-5	1	4	-5	0	6	-6	-5	-3	-3	-1	1	
18 S	1	5	-2	-4	-3	0	-1	-1	-3	-2	-3	2	-2	-5	-4	2	-1	-5	-4	0	
19 S	3	-4	-3	-3	-3	-2	-3	-1	-4	-1	0	-3	5	1	-4	2	1	-4	-1	0	

Figure 1. Sample PSSM for a sequence.

The PSSM contains 20 values for each residue in a protein sequence. For a protein sequence of N residues, there will be  $N * 20$  PSSM Values.

Now, for feature generation a sliding window of 5 around central residue is used, as shown in Figure 2.

1 T	-1	-3	-2	-3	-3	-2	-2	-3	-3	-3	-3	-2	-3	-4	-3	3	6	-4	-3	-2
2 I	-3	-5	-6	-6	-3	-5	-5	-6	-6	7	0	-5	2	-2	-5	-5	-2	-5	-4	2
3 K	-3	7	-2	-3	-5	0	-2	-4	-3	-5	-5	4	-4	-5	-4	-1	1	-5	-4	-4
4 V	-2	-5	-5	-6	-3	-5	-5	-6	-6	3	1	-5	-1	-3	-5	-4	-2	-5	-4	6
5 L	-2	-5	-6	-6	-3	-4	-5	-6	-5	1	5	-5	3	-1	-5	-4	-3	-4	-3	2
6 F	-3	-5	-6	-6	-3	-5	-5	-6	-5	5	4	-5	-1	1	-5	-5	-3	-4	-3	3
7 V	4	-4	-5	-5	1	-4	-4	-4	-5	0	-2	-4	-2	-4	-4	-3	-2	-5	-4	6
8 D	-4	-4	-1	8	-6	-2	0	-4	-3	-6	-6	-3	-6	-6	-4	-2	-3	-7	-6	-6
9 D	-4	-4	-1	8	-6	-3	-1	-4	-3	-6	-6	-3	-6	-6	-4	-2	-3	-7	-6	-6
10 H	-3	-2	-2	-3	-5	6	-1	-4	9	-5	-5	-2	-3	-4	-4	-3	-3	-5	-1	-5

**Figure 2. Sliding window of 5 around the central residue 'L'.**

For the residue L, we consider the PSSM values of K & V above it and the PSSM values of F & V below it to generate a feature vector. Therefore, there will be  $20 * 5 = 100$  PSSM values for each non-terminal residue like L. For terminal residues like T that do not have one or more neighbors on either side, rows containing 20 values of -1 are used for feature generation.

The  $N * 20$  PSSM values are now converted into  $N * 100$  Feature vectors, as shown below in Figure 3.

20 PSSM values for residue 'L':

5 L	-2	-5	-6	-6	-3	-4	-5	-6	-5	1	5	-5	3	-1	-5	-4	-3	-4	-3	2
-----	----	----	----	----	----	----	----	----	----	---	---	----	---	----	----	----	----	----	----	---

100 Feature vector for residue 'L':

5 L	-2	-5	-6	-6	-3	-4	-5	-6	-5	1	5	-5	3	-1	-5	-4	-3	-4	-3	2
	-2	-5	-5	-6	-3	-5	-5	-6	-6	3	1	-5	-1	-3	-5	-4	-2	-5	-4	6
	-3	7	-2	-3	-5	0	-2	-4	-3	-5	-5	4	-4	-5	-4	-1	1	-5	-4	-4
	-3	-5	-6	-6	-3	-5	-5	-6	-5	5	4	-5	-1	1	-5	-5	-3	-4	-3	3
	4	-4	-5	-5	1	-4	-4	-4	-5	0	-2	-4	-2	-4	-4	-3	-2	-5	-4	6

**Figure 3. Feature generation from  $N * 20$  PSSM Values.**

Now, we combine the 100 PSSM values of each pair of residues (i, j), where  $j > i + 5$ , in a protein sequence i.e. for a protein sequence of length N, we will have  $N * 200$  PSSM values.

The actual intra residue distance (d) of the two residues (i, j) from the corresponding .rr files has been used to generate the class labels as 1 if (i, j) is in contact i.e.  $d < 8$ , 0 otherwise. Also, all the residue-residue which are not in contact with a class label 0, have been considered in feature generation.

A sample training instance is shown in Figure 4. Each training instance consists of i, j, 200 PSSM values and the corresponding class label for the i, j residue pair.

A sample training instance:

```
1
33
(-0.0255238280445, -0.0255238280445, -0.0255238280445, ..., -0.153142968267, -0.127619140223, -0.0255238280445)
1
```

Figure 4. Feature generation (i, j, 200 PSSM Values, class label).

## 2.1. Logistic Regression Learning on Training Set:

Initially, we implemented the Stochastic gradient ascent based optimization algorithm and MCLE estimation to learn the weight vector of Logistic Regression using the training set.

In Stochastic Gradient Ascent, we begin by selecting a single training instance for each iteration using simple random sampling with replacement and initially all the weights are set to zero. In every iteration, we calculate a prediction based on the logistic function using the training instance and the weight vectors as follows:

$$P(z) = \frac{1}{1+e^z}$$

where,  $P(z)$  = Probability estimate

$z$  = input to the function ( $w_0 + \sum_i w_i x_i$ )

The weight vectors are updated in every iteration using the MCLE estimation as follows:

$$w_i \leftarrow w_i + \eta X_i (Y - P(Z))$$

where,

$w_i$  : Weights

$\eta$  : learning rate

$X_i$  : Features from the selected training instance

$Y$  : Prediction

$P(z)$  = Probability estimate

Assumption :  $X_0 = 1$  for  $W_0$

The learning algorithm uses gradient ascent to discover best values of regression coefficients based on the error in prediction on training data in an iterative process, until convergence is reached. Given a dataset, we want to find the parameter vector ' $\mathbf{w}$ ' which maximizes the likelihood.

During the training, we use a variable learning rate ( $\eta$ ) which follows an exponential function and varies from 0.045 to 0.009405, i.e. the learning rate is high at higher gradient value and reduces as gradient decreases and we are closer to the maxima.

The above process is repeated until the error in our prediction doesn't change over a few iterations, i.e. change in the weights is less than  $\epsilon = 0.0005$  for 10 iterations.

Also, we have implemented mini batch gradient ascent optimization algorithm in combination with MCLE estimation to learn the weight vector of Logistic Regression using the training set.

In mini batch gradient ascent, we begin by selecting a subset of the large training dataset, say, 500 training instances using random sampling with replacement and repeat the above mentioned process of learning weights.

## **2.2. Challenges faced in training:**

1. **Normalization of dataset:** We observed that, the PSSM values for the 150 protein sequences were not normalized, due to which weights learned using gradient ascent optimization were too large and non-converging with the use of general learning rates and stopping criteria.

Hence, we normalized our features vectors using Euclidean distances. Using Euclidean distances, our feature vector values were normalized in the range of 0 and 1.

2. **Batch Gradient Ascent:** The implementation of batch gradient ascent on the whole training dataset was computationally intensive and time taking. Hence, we used mini batch gradient ascent where we learn the weights on a randomly selected small subset of the training dataset.

## **2.3. Logistic Regression Classification on Test Set:**

For classification, we use the weights learned as a part of training the binary logistic model and the logistic sigmoid function to classify the residue-residue pairs in the protein sequence as 1 if in contact, 0 otherwise.

The residue-residue pairs are classified as in contact (1) if the probability of the prediction is greater 0.5, not in contact (0) otherwise.

The prediction is calculated using the following equation:

$$P(z) = \frac{1}{1+e^z}$$

where,  $P(z)$  = Probability estimate

$z$  = input to the function ( $w_0 + \sum_i w_i x_i$ )

### Accuracy:

The ‘top’ L/10, L/5, L/2 predicted contacts are used to evaluate the classification performance averaged over the proteins in the test dataset, where L is the length of the protein sequence and ‘top’ contacts are the contacts predicted with high probabilities.

### 3. Results

The following is a sample output showing the residue-residue contacts in a protein sequence in RR format. The contacts are sorted in the non-increasing their probabilities:

```

KTRWTRREEDEKLKKLVEQNGTDDWKVIANYLPNRTDVQCQHRWQKVLNPE
8 17 0 8 0.510280046958
8 50 0 8 0.510017163785
9 17 0 8 0.509908357755
9 50 0 8 0.509645466791
7 17 0 8 0.509619764289
7 50 0 8 0.509356867476
8 36 0 8 0.509345703421
8 22 0 8 0.509339780959
8 48 0 8 0.509103081724
8 32 0 8 0.509072664174
8 21 0 8 0.509068859846
9 36 0 8 0.50897398746
9 22 0 8 0.508968064837
8 29 0 8 0.508807248786
9 48 0 8 0.508731359239
9 32 0 8 0.508700940884
9 21 0 8 0.508697136455
7 36 0 8 0.508685373931
7 22 0 8 0.508679451187
8 33 0 8 0.508534607826
7 48 0 8 0.50844274083
9 29 0 8 0.508435518584
7 32 0 8 0.508412321872
7 21 0 8 0.508408517369
10 17 0 8 0.508232047148
9 33 0 8 0.508162870742
7 29 0 8 0.508146894409
23 50 0 8 0.507986955242
8 23 0 8 0.507979473289
10 50 0 8 0.507969124657
8 20 0 8 0.507962141478
7 33 0 8 0.507874241431
8 45 0 8 0.507822080078
8 41 0 8 0.507695762295
9 23 0 8 0.507607722876
9 20 0 8 0.507590390664
8 14 0 8 0.5075903665

```



## The following are the accuracies of our logistic model trained using **Stochastic Gradient Ascent**:

Protein	L10		L5		L2		TP		TN		FP		FN
1gzc.rr	82.6086956522	%	91.4893617021	%	93.2773109244	%	0		26535		0		726
1jos.rr	90.0	%	95.0	%	98.0	%	0		4296		0		169
5ptp.rr	95.4545454545	%	97.7272727273	%	98.1981981982	%	0		22783		0		653
1dbx.rr	93.3333333333	%	93.3333333333	%	97.3333333333	%	0		10053		0		387
1lo7.rr	92.8571428571	%	96.4285714286	%	98.5714285714	%	0		8758		0		287
1w0h.rr	94.7368421053	%	97.3684210526	%	98.9690721649	%	0		17363		0		403
1aba.rr	87.5	%	94.1176470588	%	97.6744186047	%	0		3181		0		140
1kw4.rr	83.3333333333	%	92.3076923077	%	97.0588235294	%	0		1881		0		72
1h4x.rr	90.9090909091	%	90.9090909091	%	94.5454545455	%	0		5265		0		195
1l1j.rr	90.0	%	95.2380952381	%	98.1132075472	%	0		4795		0		255
1hfc.rr	93.3333333333	%	96.7741935484	%	97.4358974359	%	0		11137		0		339
1tlf.rr	85.7142857143	%	93.3333333333	%	94.7368421053	%	0		2369		0		116
1g2r.rr	88.8888888889	%	88.8888888889	%	95.7446808511	%	0		3757		0		159
2mhr.rr	90.9090909091	%	95.652173913	%	96.6101694915	%	0		6179		0		149
1pch.rr	87.5	%	94.1176470588	%	95.4545454545	%	0		3218		0		185
1bdo.rr	75.0	%	87.5	%	90.0	%	0		2583		0		192
2cua.rr	91.6666666667	%	95.8333333333	%	98.3686557377	%	0		6473		0		313
1c52.rr	84.6153846154	%	92.3076923077	%	96.9230769231	%	0		7639		0		236
1aoe.rr	94.7368421053	%	97.3684210526	%	98.9583333333	%	0		16988		0		403
1ag6.rr	88.8888888889	%	94.7368421053	%	97.9591836735	%	0		4134		0		237
1ku3.rr	83.3333333333	%	91.6666666667	%	93.3333333333	%	0		1483		0		57
1c3w.rr	81.25	%	84.8484848485	%	86.7469879518	%	0		12531		0		349
1a3a.rr	92.8571428571	%	96.5517241379	%	98.6111111111	%	0		9405		0		325
1mk0.rr	77.7777777778	%	89.4736842105	%	91.6666666667	%	0		4005		0		181
1ek0.rr	93.75	%	96.9696969697	%	98.8095238095	%	0		12807		0		396
1i4j.rr	90.9090909091	%	95.4545454545	%	98.1818181818	%	0		5230		0		230
1l5g.rr	92.8571428571	%	92.8571428571	%	95.8333333333	%	0		9291		0		300
1xdz.rr	95.652173913	%	97.8723404255	%	98.3193277311	%	0		26483		0		545
1dgg.rr	84.6153846154	%	88.4615384615	%	91.0447761194	%	0		7910		0		346
2arc.rr	93.75	%	96.875	%	98.75	%	0		11742		0		348
1fgt.rr	90.0	%	85.7142857143	%	90.7407407407	%	0		5075		0		281
1jbk.rr	94.4444444444	%	97.2972972973	%	96.8085106383	%	0		16481		0		355
1a70.rr	88.8888888889	%	94.7368421053	%	97.9166666667	%	0		3966		0		220
1l58.rr	94.4444444444	%	97.2972972973	%	98.9361702128	%	0		16476		0		360
1g90.rr	77.7777777778	%	83.3333333333	%	88.8888888889	%	0		3472		0		183
1lin.rr	95.4545454545	%	97.7272727273	%	99.1071428571	%	0		23301		0		570
1brf.rr	80.0	%	90.0	%	84.6153846154	%	0		1037		0		91
1atl.rr	95.0	%	97.5	%	99.0	%	0		18462		0		453
Average L/10 = 89.0723292642 %													
Average L/5 = 93.5544516791 %													
Average L/2 = 95.8219740864 %													
Average accuracy positive contacts = 0 %													
Average accuracy negative contacts = 96.9693035835 %													

We achieve an **accuracy of 0% in stochastic gradient ascent** because our model predicts all the probabilities as 0's. This is due to the fact that the weight bias  $W_0$ , dominates all the other 200 weights trained by our model. The following are the values of the weights learned using stochastic gradient ascent:

```

===== recent 10 deltas =====
[1.1822106803050405e-11, 1.642955435095226e-11, 4.284011873530376e-11, 4.7769105653930308e-11, 4.871994928761673e-11, 1.6294877444644457e-11, 3.124041492582575e-11, 3.599949094557515e-11, 4.068426132521334e-11, 7.662119415931481e-11]

===== weights =====
[[4.862120852093954, -0.3068340424467362, -0.5645102324261082, -0.5919559154002027, -0.6359521475922265, -0.8748815468092239, -0.4930407319065049, -0.5192237281350535, -0.6559887144139089, -0.61351842392457
99, -0.5008423828157083, -0.5893203095213102, -0.47421024502506327, -0.5062396524322125, -0.72258581193087782, -0.8418707566299786, -0.4067223814493749, -0.3887836292866165, -1.1022847140919125, -0.703765894
6550015, -0.38782066586006514, -0.2673175014104722, -0.6081460744947806, -0.726029514299328, -0.7701833908239132, -0.7844040343754791, -0.5674813421965329, -0.6592163310764798, -0.782373449386169, -0.67560
08830804719, -0.38810283709551325, -0.49573283156700765, -0.5565480819552907, -0.4233849129404532, -0.866249586606149, -0.942730470345319, -0.44457529384206433, -0.3395950090799228, -1.0289972891813013, -0
.6775480823792572, -0.2345973516195558, -0.2971783680596982, -0.8202607171059093, -0.8419641149305267, -0.9571630252278982, -0.7106034838502338, -0.7275384927400609, -0.8308606933011238, -0.8313687918335759
, -0.7808809997852897, -0.24865811386406916, -0.32046043592714624, -0.7750738247996082, -0.37792949860018404, -0.5268967874788394, -0.9838331550393444, -0.5382814518071383, -0.40888660292559964, -1.04346560
80421754, -0.6383827883754704, -0.14296756621174211, -0.29955405587812983, -0.6357747003049424, -0.658293127991988, -0.7056795720337099, -0.8623305508327862, -0.5360038479949217, -0.5702834623984521, -0.8112
868939519353, -0.6982666127062951, -0.45951235989296424, -0.5644537834243267, -0.5729699179522403, -0.4934420229057796, -0.684052383234774, -0.9308190289102966, -0.43464555588488873, -0.36374657173909397,
-1.058444608522129, -0.658876917953763, -0.3209963641638313, -0.27254018201379408, -0.53780703675909478, -0.5864056080441187, -0.6381152991298793, -0.8911546351416636, -0.43835708520997056, -0.480397695014773
14, -0.652258886404227, -0.579370385899744, -0.6270229894223799, -0.6553513221862953, -0.43171975400025875, -0.5527034201451749, -0.7612109075704906, -0.8715080557530678, -0.37045135973126847, -0.3764580570
7895551, -1.0266823870434748, -0.6477946573351032, -0.5401732888750778, -0.22935157981328855, -0.4616355242059634, -0.52494073112021, -0.5186123352469275, -0.851482757921263, -0.403464522628266, -0.374858
11248634715, -0.5689470643496338, -0.5524490695347273, -0.6026261077040717, -0.632293033114406, -0.36927585116671724, -0.5206724295636918, -0.819106086846733, -0.8601078114164373, -0.38846414899706544, -0.3
7080558029451216, -1.1052898489781595, -0.6782881331165902, -0.4230513615614961, -0.2655160652825049, -0.4549173258861444, -0.5328051732449073, -0.559455418112547, -0.8933479746895162, -0.41824476310192416
, -0.439336903708605, -0.7394503204376078, -0.5726206122446678, -0.4878254376395328, -0.5412159530039465, -0.4012026308896511, -0.5091434567054046, -0.7006775692842923, -0.8051664131632604, -0.390021957585
9059, -0.3275434512708515, -0.998391559878311, -0.5997998555174555, -0.3539545353238082, -0.2907433797572786, -0.7202704133415879, -0.8016494779228029, -0.9045629211139206, -0.7524811502714239, -0.624026804
4696639, -0.7398345845164525, -0.8517072885428768, -0.7428198322940265, -0.2470007821143372, -0.34659313103375066, -0.6768083630881659, -0.37164790146589183, -0.5746864563552234, -0.8757203789313087, -0.553
1760684414949, -0.42990089397329665, -1.074797290649172, -0.6326278010709231, -0.17466282511730544, -0.322467172900829, -0.5455829994134067, -0.6244041197361697, -0.6801017685153092, -0.8533423867813058, -
0.4887570633403971, -0.5413285621046004, -0.8045174102918962, -0.5821368955369175, -0.4170819589112943, -0.4776618509231465, -0.48170448958041927, -0.4812423178598619, -0.649461295994092, -0.847928911841034
8, -0.46195745159873736, -0.3412058266012254, -1.0258541407277294, -0.6127633521247882, -0.2765858429895548, -0.24313758932377916, -0.48924094211931457, -0.4989886611584737, -0.49196347432096, -0.9664212764
498331, -0.39230155155475627, -0.4041978919394768, -0.6187359622471781, -0.601739510463288, -0.597214738724545, -0.6516094742849202, -0.3838024658387346, -0.5649168608074479, -0.8627911391225757, -0.684882
2856379218, -0.3666018421248239, -0.3996698383437451, -1.1263752770859805, -0.726085039110758, -0.48267986513611094]

```



The value of  $W_0$  is 14.06, whereas all the other weights are negative values.

The following are the accuracies of our logistic model trained using **mini Batch Gradient Ascent**:

Protein	L10		L5		L2		TP		TN		FP		FN
1gzc.rr	0.0	%	0.0	%	0.0	%	123		18401		8134		603
1jos.rr	0.0	%	0.0	%	0.0	%	17		2893		1403		152
5ptp.rr	0.0	%	0.0	%	0.0	%	84		17166		5617		569
1ddx.rr	0.0	%	0.0	%	0.0	%	43		7340		2713		344
1l07.rr	0.0	%	0.0	%	0.0	%	25		7009		1749		262
1w0h.rr	0.0	%	0.0	%	1.03092783505	%	62		12179		5184		341
1aba.rr	0.0	%	0.0	%	0.0	%	11		2130		1051		129
1kw4.rr	0.0	%	0.0	%	0.0	%	10		1358		523		62
1h4x.rr	9.09090909091	%	4.54545454545	%	1.01818181818	%	28		3642		1623		167
1i1j.rr	0.0	%	0.0	%	0.0	%	59		2903		1892		196
1hfc.rr	0.0	%	0.0	%	1.28205128205	%	57		7751		3386		282
1t1f.rr	0.0	%	0.0	%	0.0	%	14		1291		1078		102
1g2r.rr	0.0	%	0.0	%	0.0	%	28		2129		1628		131
2mhr.rr	9.09090909091	%	8.69565217391	%	3.38983050847	%	32		4020		2159		117
1pch.rr	0.0	%	0.0	%	0.0	%	39		2114		1104		146
1bdo.rr	0.0	%	0.0	%	0.0	%	31		1826		757		161
2cua.rr	0.0	%	0.0	%	0.0	%	60		4308		2165		253
1c52.rr	0.0	%	0.0	%	0.0	%	79		4190		3449		157
1aoe.rr	0.0	%	0.0	%	0.0	%	81		11403		5585		322
1ag6.rr	0.0	%	10.5263157895	%	8.16326530612	%	43		2694		1440		194
1ku3.rr	16.6666666667	%	8.33333333333	%	6.66666666667	%	15		822		661		42
1c1w.rr	0.0	%	0.0	%	0.0	%	47		9160		3371		302
1a3a.rr	0.0	%	0.0	%	1.38888888889	%	46		5941		3464		279
1mk0.rr	0.0	%	0.0	%	0.0	%	21		2804		1201		160
1ek0.rr	0.0	%	0.0	%	0.0	%	45		9123		3684		351
1l4j.rr	0.0	%	9.09090909091	%	7.27272727273	%	40		3508		1722		190
1l5g.rr	0.0	%	0.0	%	0.0	%	46		6073		3218		254
1xdz.rr	0.0	%	0.0	%	0.0	%	79		19530		6953		466
1dgg.rr	0.0	%	0.0	%	0.0	%	83		4299		3611		203
2arc.rr	0.0	%	0.0	%	0.0	%	68		8445		3297		280
1fgt.rr	0.0	%	0.0	%	1.85185185185	%	40		3754		1321		241
1j1bk.rr	5.55555555556	%	2.7027027027	%	1.06362978723	%	74		10201		6200		281
1a70.rr	0.0	%	0.0	%	0.0	%	34		2735		1231		186
1l58.rr	0.0	%	0.0	%	0.0	%	81		10009		5567		279
1g90.rr	0.0	%	0.0	%	0.0	%	29		2091		1381		154
1i1n.rr	0.0	%	0.0	%	0.892857142857	%	92		16499		6802		478
1brf.rr	0.0	%	0.0	%	0.0	%	15		762		275		76
1atl.rr	0.0	%	0.0	%	0.0	%	44		14056		4406		409
Average L/10 = 1.06326422116 %													
Average L/5 = 1.15511493778 %													
Average L/2 = 0.916344167371 %													
Average accuracy positive contacts = 1.61633159153 %													
Average accuracy negative contacts = 96.3475315371 %													

We achieve an **accuracy of 1.06% in mini Batch gradient ascent**. This is due to the fact that the weight bias  $W_0$ , dominates all the other 200 weights trained by our model. Due to the dominating  $W_0$  value, our model is forced to predict more number of False Positives and less number of True Positives. The following are the values of the weights learned using stochastic gradient ascent:

```
===== recent 10 deltas =====
[0.0004962227714559144, 0.00029811988686411205, 0.0002987962735179193, 0.0001005263567379975, 0.0003997944024781512, 0.0003993377492147533, 0.00019976385561127214, 0.0002008596041590747, 0.0002004433654922637, 0.00029858138599702825]

===== weights =====
[-0.00041416718032910794, 2.5155895854366857e-05, -0.0002296455335513812, -0.00023245256121637124, -0.00013109929224467982, 0.00016340607253262653, -0.00013400985343908932, -0.00013713375695845252, -5.9467149377576455e-05, 1.9092613158317915e-05, 0.00024471208617760445, -5.823227237711335e-05, -0.0001692606807960826, -3.935695625902801e-05, 4.980979357927851e-05, -6.982776468458289e-05, -4.861433882397066e-05, -1.625734023406714e-05, -9.748038321925846e-05, -9.87437689855272e-06, 0.0002465298568977053, -6.126983694134217e-05, -0.00020097470286527337, -0.00047696963125386596, -0.0006018571181277125, 0.0003924593692384824, -0.00047716134223771975, -0.0005512073524640802, -0.0002509948886822563, -0.000349284793172768, 0.00040579050758909915, 0.00018926103477553301, -0.0003707923947783986, 0.00017916884444498304, 0.00012800537378964723, 1.9128151672717846e-05, -0.0002717288040855121, -0.00011008526324435759, 0.00019254983127692805, -1.2514564324737794e-05, 0.00040547025654078135, -0.00025288141630022833, -0.0009546360002935029, -0.000919788848697887, -0.0011834792382001242, 0.000693428473953401, -0.0009823700079957674, -0.0013280043583451656, -0.0005738481078271662, -0.0005328291802301311, 0.001041686975552783, 0.0006217038274499545, -0.0010820381338226977, 0.00040343401758084627, 0.0009073396569629247, -0.0005385034649523752, -0.0005826832617877308, -0.0002943233241214008, 0.0004940455039870742, 0.0004985608321752142, 0.0008573255780513987, -7.778189236161496e-05, -0.00012160282938356424, -0.0004580728681829073, -0.0004009919495683538, 0.00024844017718793365, -0.0002561765720355885, -0.00016808286887567573, -0.0004579581046839585, -0.0003772908730425127, 0.00048648072647170106, 0.00040236646543271177, -0.00015882456719607264, 0.0002507164305061206, 0.000531207929155855, -3.7454783162185464e-05, -0.00022643157672091896, 0.00014117793582598515, 0.0003653131461790603, 0.00038572122942463527, 0.0004602502841304144, -0.00020908346837666787, -2.9213134259182635e-05, 0.0001706601533221452, 0.00016136257217664555, 0.00017942990361628986, -5.177880379812575e-05, 0.000648143720143e-05, 0.0002562977353410643, 0.0002521264213348209, -0.0001680847662615145, -0.0001587345893219544, -0.00016418792694208805, -0.0001298948249671039, 0.00014478786403149951, -8.191234398355285e-05, 8.257907368359922e-05, -0.00013592545779585354, -5.213524881972107e-07, 0.00013323393738511676, -0.00016657881972816015, -0.0001740607740154289, -0.0001889989537102009, -0.00013622029929490363, 5.455768356790573e-05, 0.00013868733217433626, -7.846527632168012e-05, 0.0001409751505928046, 2.6402448080735014e-05, -3.222194054768465e-05, 0.00023206772630286523, 4.4118688390188376e-05, 7.881892342764353e-05, -8.532438259544798e-05, 0.00010179274205966326, 1.692290513902842e-05, -0.0001561147608412354, -2.2763830577324896e-05, -1.0006765166621265e-05, 1.7514748308965183e-05, 0.000399116311761613, -0.00017655119482680394, -6.403166162003415e-05, -0.00014250582676922005, -0.0002917692109689933, 0.0003950975199273933, -0.00011814613148507777, -5.849119863168329e-05, -0.00023988805266381076, 0.00011420699319116188, 0.00042180206845820025, 0.00028188979626472, 474, -0.0001837402556595982, 0.00015838790712894408, 0.00018340425779848337, -0.00018021887100111363, -0.00012781138421335767, 0.0001801007660828135, 4.155167614930432e-05, 0.00012189581203242815, 0.0004980498034059674, -0.0001585932359863328, -0.0009934995322410291, -0.00106088522356255, -0.001556423144302265, 0.000779427005910186, -0.0010073737970618338, -0.0014098434346578263, -0.0005964993527854978, -0.00044455096816455213, 0.0001343057955736012, 0.00010856445318117316, -0.0011685473430408813, 0.00006639338095727135, 0.0008725562615835435, -0.0006365882266444687, -0.000532785736948616, -7.878590742940436e-05, 0.0004721226392785114, 0.0003833106711498594, 0.0009789908415117458, -0.00025670552298333955, -0.0001922785914651942, -9.058075693677188e-05, -0.0003230971121202341, 0.000431978388793443, -0.0002925452609283615, -0.000328527385171846, -0.0002608991563883416, -8.460883917648554e-05, 0.0002879989526956179, 6.462632065979517e-05, -0.0002518613440901785, -2.571411577374214e-05, 0.000202623628877077315, -0.00022766279394939324, -0.00015915932818219295, 5.22894470559667e-05, 0.00016086285843326545, 0.0001643118550066359, 0.0003298112532174322, -0.0002669787882153092, -0.0001181958820402733, 5.533792353966006e-05, 0.00020512558410691319, 0.000294493923741761, -5.0743878521652254e-05, -1.634031092630102e-05, -2.7389887487770786e-05, 5.030162536075467e-05, -3.3039074963952085e-05, -0.00019778566001601785, -0.000273644352000947, -0.0001846515801735244, 3.803141941583929e-05, -0.00030264950171965, -5.765801800876261e-05, -0.00013862613325452757, 0.00013014011149186923, -4.698710996002354e-05, -1.7206180711757466e-05]
```

## 4. Challenges due to Class Imbalance

We observed that, the number of non-contact pairs in the training dataset largely outnumbered the number of contact pairs. Due to this class imbalance, initially our model was incorrectly trained to predict non-contact pairs.

Hence, we balanced our dataset with equal number of contact and non-contact pairs. We randomly selected non-contact pairs to equate them to the number of contact pairs.

Before class balancing, our overall prediction accuracy was 97% and both of our models (stochastic & batch) were only predicting non-contact pairs. Due to class balancing, both of our models improved, and the model trained using mini batch gradient ascent gained a contact accuracy of 1.06%.

## 5. Discussion

We performed different experiments as discussed above to improve the accuracy of our binary logistic model. Our contact accuracy was the 0 % and our model was only predicting non-contacts due to class imbalance. When the contact and non-contact pairs in our training dataset were balanced, our accuracy increased to 1.06 %. These experiments outlay the importance of class imbalance, feature normalization and mini batch gradients.

## 5. References

1. Wikipedia contributors. (2018, March 21). Protein structure. In Wikipedia, The Free Encyclopedia. Retrieved 19:23, March 26, 2018, from [https://en.wikipedia.org/w/index.php?title=Protein\\_structure&oldid=831674619](https://en.wikipedia.org/w/index.php?title=Protein_structure&oldid=831674619)
2. Pauling, L., Corey, R. B., & Branson, H. R. (1951). The structure of proteins: two hydrogen-bonded helical configurations of the polypeptide chain. *Proceedings of the National Academy of Sciences*, 37(4), 205-211.
3. Alberts, B., Johnson, A., Lewis, J., Raff, M., Roberts, K., & Walter, P. (2002). The shape and structure of proteins.
4. Anfinsen, C. B. (1973). Principles that govern the folding of protein chains. *Science*, 181(4096), 223-230.
5. Zhang, H. (2004). The optimality of naive Bayes. *AA*, 1(2), 3.
6. Schmidler, S. C., Liu, J. S., & Brutlag, D. L. (2000). Bayesian

- segmentation of protein secondary structure. *Journal of computational biology*, 7(1-2), 233-248.
7. Robles, V., Larrañaga, P., Peña, J. M., Menasalvas, E., Pérez, M. S., Herves, V., & Wasilewska, A. (2004). Bayesian network multi-classifiers for protein secondary structure prediction. *Artificial Intelligence in Medicine*, 31(2), 117-136.
  8. Di Lena, P., Nagata, K., & Baldi, P. (2012). Deep architectures for protein contact map prediction. *Bioinformatics*, 28(19), 2449-2457.
  9. Tress, M. L., & Valencia, A. (2010). Predicted residue–residue contacts can help the scoring of 3D models. *Proteins: Structure, Function, and Bioinformatics*, 78(8), 1980-1991.